

regression_task

April 3, 2019

1 Data discription

1. **FIRMCOST** - The measure of the firm's risk management cost effectiveness, defined as total property and casualty premiums and uninsured losses as a percentage of total assets.
2. **ASSUME** - Per occurrence retention amount as a percentage of total assets.
3. **SIZELOG** - Logarithm of total assets.
4. **INDCOST** - A measure of the firm's industry risk.
5. **CENTRAL** - A measure of the importance of the local managers in choosing the amount of risk to be retained.
6. **SOPH** - A measure of the degree of importance in using analytical tools.

```
In [3]: import pandas as pd
import numpy as np
```

2 Display data statistics

```
In [4]: def corr(df, col_names):
        '''
            compute correlation matrix
        '''
        rs = np.random.RandomState(0)
        df = pd.DataFrame(rs.rand(len(col_names), len(col_names)), columns=col_names)
        corr = df.corr()
        corr = corr.style.background_gradient(cmap='coolwarm').set_precision(2)

        return corr

In [5]: # define path to RiskSurvey data
PATH_TO_DATA = 'regression-data/data-actuarial/RiskSurvey.csv'

# creating dataframe
df = pd.read_csv(PATH_TO_DATA)
df = df.drop('CAP', axis=1)
```

```
In [6]: # display data
df.head()
```

```
Out [6]:
```

	FIRMCOST	ASSUME	SIZELOG	INDCOST	CENTRAL	SOPH
0	3.29	0.29	9.55	0.32	1	25
1	9.31	0.89	8.04	0.33	2	24
2	4.07	1.67	7.90	0.34	2	15
3	6.94	1.21	8.10	0.34	1	16
4	5.35	0.28	7.74	0.09	3	18

```
In [7]: # data description
df.describe()
```

```
Out [7]:
```

	FIRMCOST	ASSUME	SIZELOG	INDCOST	CENTRAL	SOPH
count	73.000000	73.000000	73.000000	73.000000	73.000000	73.000000
mean	10.973288	2.573562	8.331918	0.418356	2.246575	21.191781
std	16.158611	8.444978	0.963378	0.216243	1.255884	5.303713
min	0.200000	0.000000	5.270000	0.090000	1.000000	5.000000
25%	3.510000	0.240000	7.650000	0.330000	1.000000	18.000000
50%	6.080000	0.510000	8.270000	0.340000	2.000000	23.000000
75%	12.710000	1.670000	8.950000	0.500000	3.000000	25.000000
max	97.550003	61.820000	10.600000	1.220000	5.000000	31.000000

3 Show correlation between predictors

```
In [8]: # build correlation matrix of predictors
corr(df.loc[:, 'FIRMCOST':'SOPH'], ['FIRMCOST', 'ASSUME', 'SIZELOG', 'INDCOST', 'CENTRAL'])
```

```
Out [8]: <pandas.io.formats.style.Styler at 0x7f2a8000b208>
```

4 Fit linear model

```
In [9]: def fit_linear_model(df):
    input_matrix = df.copy()
    input_matrix['INTERCEPT'] = pd.Series(1, index=input_matrix.index)

    X = input_matrix.loc[:, 'ASSUME':'INTERCEPT'].values
    y = input_matrix['FIRMCOST'].values

    coefs = np.dot(np.dot(np.linalg.inv(np.dot(X.transpose(), X)),
                                X.transpose()), y)

    return coefs

# split dataset on train and test 60:40, due to low number of examples
X_train, X_test = df[:40], df[40:]
# compute coefficient for linear model
coefs = fit_linear_model(X_train)
coefs
```

```
Out[9]: array([ 1.65692943e-01, -2.06567267e+00,  3.21444851e+01, -2.46195547e-02,
               -1.91112242e-01,  1.82712436e+01])
```

5 Get predictions

```
In [10]: # get results on test dataset
def linear_model_predict(coefficients, to_predict):
    pred = to_predict.apply(
        lambda x: sum(list(map(
            lambda k: k[0]*k[1], zip(np.append(x.values[1:], 1), coefficients))))
        , axis=1)

    return pred

# comparison
y_predictions = linear_model_predict(coefs, X_test)
comp = y_predictions.to_frame(name='predictions')
comp['actual'] = X_test['FIRMCOST']
comp['residual'] = comp['actual'] - comp['predictions']
comp
```

```
Out[10]:
```


	predictions	actual	residual
40	7.160693	2.160000	-5.000693
41	4.355577	0.360000	-3.995577
42	5.532433	7.830000	2.297567
43	9.434665	5.090000	-4.344665
44	-4.519594	0.200000	4.719594
45	7.408182	8.850000	1.441818
46	9.943768	0.760000	-9.183768
47	5.365878	5.710000	0.344122
48	27.372334	17.530001	-9.842334
49	5.918854	14.000000	8.081146
50	5.524667	2.060000	-3.464667
51	4.203710	0.930000	-3.273710
52	12.420589	10.000000	-2.420589
53	14.960786	5.820000	-9.140786
54	5.799707	9.130000	3.330293
55	14.295559	9.000000	-5.295559
56	28.505357	12.610000	-15.895358
57	5.473248	2.150000	-3.323248
58	9.991330	22.219999	12.228669
59	9.269259	12.710000	3.440741
60	14.418950	15.970000	1.551050
61	5.914935	4.320000	-1.594934
62	7.990755	8.490000	0.499245
63	13.591023	5.250000	-8.341023
64	23.189775	18.330000	-4.859775

65	12.571876	21.719999	9.148123
66	5.913659	0.400000	-5.513659
67	9.121543	3.700000	-5.421543
68	11.491456	15.000000	3.508544
69	23.757871	18.000000	-5.757871
70	8.911315	29.120001	20.208686
71	13.261827	79.300003	66.038176
72	33.217764	13.570000	-19.647764

6 Residual plot

```
In [11]: import matplotlib.pyplot as plt
```

```
plt.plot(comp['predictions'], comp['residual'], 'ro', np.linspace(0, max(comp['predictions'], 35), 100))
plt.xlabel('predicted')
plt.ylabel('residuals')
plt.show()
```




There is no visible pattern but we have several outliers.

```
In [12]: # density plot
import seaborn as sns
sns.distplot(comp['residual'], hist=False)
```

```
/home/daniil/anaconda3/lib/python3.6/site-packages/scipy/stats/stats.py:1713: FutureWarning: U
return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```

Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x7f2a68844ac8>



7 Compute R^2

```
In [13]: def r_squared(actual, predicted):
    mean = sum(actual)/len(actual)
    RSS = sum(list(map(lambda x: (x[0] - x[1])**2, zip(predicted, actual))))
    TSS = sum(list(map(lambda x: (x - mean)**2, actual)))

    return 1 - RSS/TSS

score = r_squared(comp['actual'].tolist(), comp['predictions'].tolist())
score
```

Out[13]: 0.016972853550193512

R^2 of the model equals to **0.016**, so the model fits a bit better than a horizontal hyperplane. Due to high correlation between predictors and rank of X is less than $p + 1$ ($5 < 6$), so we can assume that ridge regression can be a better choice for our problem.

8 Estimating covariance matrix of \tilde{y}

```
In [14]: # Variance of
score = r_squared(comp['actual'].tolist(), comp['predictions'].tolist())
X = X_train.copy()
X['INTERCEPT'] = pd.Series(1, index=X.index)
y = X['FIRMCOST'].values
X = X.loc[:, 'ASSUME':'INTERCEPT'].values
H = np.dot(np.dot(X, np.linalg.inv(np.dot(X.transpose(), X))), X.transpose())
# estimating  $\tilde{\sigma}^2$ 
sigma2 = np.dot(np.dot(y.transpose(), (np.identity(H.shape[0]) - H)), y)/(X.shape[0] - H.shape[0])
var = np.dot(np.linalg.inv(np.dot(X.transpose(), X)), sigma2)
pd.DataFrame(var)
```

Out[14]:

	0	1	2	3	4	5
0	1.264688	0.530214	-2.308232	0.044273	-0.052752	-4.163613
1	0.530214	12.687781	0.411645	1.493031	0.760806	-125.394404
2	-2.308232	0.411645	150.220666	5.600334	-1.571776	-41.155763
3	0.044273	1.493031	5.600334	6.343226	-0.295370	-22.117418
4	-0.052752	0.760806	-1.571776	-0.295370	0.333641	-11.916516
5	-4.163613	-125.394404	-41.155763	-22.117418	-11.916516	1367.329784

9 Ridge regression

```
In [15]: from sklearn.preprocessing import StandardScaler

def fit_ridge_regression(x, y, alpha=1):
    X = x.copy()
    X = np.array([np.append(step, 1) for step in X])

    return np.dot(np.dot(np.linalg.inv((np.dot(X.transpose(), X) +
                                         np.dot(alpha, np.identity(X.shape[1])))),
                                         X.transpose()), y)

def ridge_reg_predict(coefficients, to_predict):
    pred = to_predict.apply(
        lambda x: sum(list(map(
            lambda k: k[0]*k[1], zip(np.append(x.values, 1), coefficients))))
        , axis=1)

    return pred

# Standarize data
scaler = StandardScaler()
X_train_rg, y_train = X_train.loc[:, 'ASSUME':'SOPH'].values, X_train['FIRMCOST'].values
X_train_std = scaler.fit_transform(X_train_rg)

# getting most suitable alpha
```

```

alphas, output = np.linspace(0, 10, 1000), []
for alpha in alphas:
    ridge_coefs = fit_ridge_regression(X_train_std, y_train, alpha=alpha)
    ridge_pred = ridge_reg_predict(ridge_coefs, pd.DataFrame(X_train_std))
    output.append(r_squared(y_train, ridge_pred))

best_alpha = alphas[output.index(max(output))]
ridge_coefs = fit_ridge_regression(X_train_std, y_train, alpha=best_alpha)
ridge_pred = ridge_reg_predict(ridge_coefs, pd.DataFrame(scaler.transform(X_test.loc[
y_test = X_test['FIRMCOST'].values
r_squared(y_test, ridge_pred)

```

Out[15]: 0.016972853550185962

R^2 score of ridge regression is the same as R^2 of pure linear model. Due to most suitable alpha, which is 0.0. So, ridge regression doesn't improve performance.

10 Applying Principal Component Analysis to reduce dimension of the space

In [16]: `import plotly.plotly as py`

```

scaler = StandardScaler()
x_train, y_train = X_train.loc[:, 'ASSUME':'SOPH'].values, X_train['FIRMCOST'].values
x_test, y_test = X_test.loc[:, 'ASSUME':'SOPH'].values, X_test['FIRMCOST'].values
scaler.fit(x_train)

# rescale our data
x_train = scaler.transform(x_train)
x_test = scaler.transform(x_test)

# getting mean vector
mean_vec = np.mean(x_train, axis=0)
# build covariance matrix
cov_mat = np.dot((x_train - mean_vec).transpose(), (x_train - mean_vec))/(x_train.shape[0]-1)
# getting eigen values and corresponding vectors
eigen_values, eigen_vecs = np.linalg.eig(cov_mat)
eigen_pairs = [(np.abs(eigen_values[i]), eigen_vecs[:,i])] for i in range(len(eigen_values))
eigen_pairs.sort(key=lambda x: x[0], reverse=True)

# compute variance explained by each principal component
var_exp = [val/sum(eigen_values)*100 for val in sorted(eigen_values, reverse=True)]
# cumulative sum of the elements
cum_var_exp = np.cumsum(var_exp)
# creat plot
trace1 = dict(
    type='bar',

```

```

        x=['PC{}'.format(i) for i in range(1, len(eigen_values) + 1)],
        y=var_exp,
        name='Individual'
    )

    trace2 = dict(
        type='scatter',
        x=['PC{}'.format(i) for i in range(1, len(eigen_values) + 1)],
        y=cum_var_exp,
        name='Cumulative'
    )

    data = [trace1, trace2]

    layout=dict(
        title='Explained variance by different principal components',
        yaxis=dict(
            title='Explained variance in percent'
        ),
        annotations=list([
            dict(
                x=1.16,
                y=1.05,
                xref='paper',
                yref='paper',
                text='Explained Variance',
                showarrow=False,
            )
        ])
    )

    fig = dict(data=data, layout=layout)
    py.iplot(fig, filename='selecting-principal-components')

```

High five! You successfully sent some data to your account on plotly. View your plot in your browser.

Out[16]: <plotly.tools.PlotlyDisplay object>

11 Projection onto the new Feature Space

In [24]: *# reducing the 5-dimensional feature space to a 4-dimensional feature subspace*
constructing of the projection matrix

```

def choose_n_comp(N):
    return np.hstack((eigen_pairs[i][1].reshape(len(eigen_values),1) for i in range(N)))

num_of_comp = 4
matrix_w = choose_n_comp(num_of_comp)

```



```

# transform our samples
x_train_tr = np.dot(x_train, matrix_w)
x_test_tr = np.dot(x_test, matrix_w)

# train linear model
coefs = fit_ride_regression(x_train_tr, y_train, alpha=0)
pred = ridge_reg_predict(coefs, pd.DataFrame(x_test_tr))
r_squared(y_test, pred)

```

/home/daniil/.local/lib/python3.6/site-packages/ipykernel_launcher.py:4: FutureWarning:

arrays to stack must be passed as a "sequence" type such as list or tuple. Support for non-seq

Out[24]: -0.24806993048618198

1 component, 36% of variance explained, R^2 on linear model = 0.027 2 components, 60% of variance explained, R^2 on linear model = -1.190 3 components, 76% of variance explained, R^2 on linear model = -0.091 4 components, 90% of variance explained, R^2 on linear model = -0.248 5 components, 100% of variance explained, R^2 on linear model = 0.0169

So, optimal number of components is still 5. Better performace result with only one component we can consider as luck due to lack of samples in our dataset and impossibility to obtain balanced decision.

12 Variable Selection

12.0.1 CRITERION-BASED PROCEDURE (Adjusted R^2)

In [56]: import itertools

```

# adjusted r2
def adj_r_squared(actual, predicted, exp_var):
    r2 = r_squared(actual, predicted)
    return 1 - (1 - r2)*(len(actual) - 1)/(len(actual) - exp_var - 1)

combinations = []

#Looping over k = 1 to k = 5 components in X
for k in range(2, x_train.shape[1] + 1):
    #Looping over all possible combinations
    for indices in itertools.combinations(range(len(x_train.transpose())), k):
        combo = np.array([x_train.transpose()[i] for i in indices]).transpose()
        reg_coefs = fit_ride_regression(combo, y_train, alpha=0)
        pred = ridge_reg_predict(reg_coefs, pd.DataFrame(np.array([x_test.transpose()
        combinations.append([adj_r_squared(y_test, pred, k), indices])

# sort combinations list by adjusted r2 in descending order

```

```
combinations = sorted(combinations, key=lambda x: x[0], reverse=True)
combinations
```

```
Out [56]: [[0.15736845008248512, (1, 3)],
[0.13107789420863258, (1, 4)],
[0.08440232852963192, (1, 3, 4)],
[0.005256447227352745, (1, 2)],
[-0.007745722688474066, (1, 2, 4)],
[-0.02716435251750382, (1, 2, 3)],
[-0.043712671980775, (1, 2, 3, 4)],
[-0.08724039547125595, (0, 1, 2)],
[-0.11427456702110295, (3, 4)],
[-0.12351984229840496, (0, 1, 2, 4)],
[-0.12465931938912012, (0, 1, 2, 3)],
[-0.14968273946623256, (2, 3)],
[-0.15524864967324192, (2, 4)],
[-0.16506921060718693, (0, 1, 2, 3, 4)],
[-0.1926115478503685, (2, 3, 4)],
[-0.24696417611594135, (0, 2)],
[-0.2816119797516208, (0, 1)],
[-0.28567329009918163, (0, 2, 3)],
[-0.3057600918280441, (0, 2, 4)],
[-0.3205340505498979, (0, 1, 4)],
[-0.33569006539354174, (0, 1, 3)],
[-0.34885986079418907, (0, 2, 3, 4)],
[-0.371973947569995, (0, 1, 3, 4)],
[-0.5136209732785726, (0, 4)],
[-0.5912135704717056, (0, 3, 4)],
[-0.6474805992304404, (0, 3)]]
```

As we can see, best adjusted R^2 score we get when using **SIZELOG** and **CENTRAL** components.

```
In [71]: best = combinations[0][1]
reg_coefs = fit_ridge_regression(np.array([x_train[:, i] for i in best]).transpose(),
pred = ridge_reg_predict(reg_coefs, pd.DataFrame(np.array([x_test[:, i] for i in best]),
r_squared(y_test, pred)
```

```
Out [71]: 0.2100329219523298
```

R^2 score = **0.21** for **SIZELOG** and **CENTRAL** features.

13 Perform F-test

```
In [83]: def f_test(full, reduced, actual, dff, dfr):
SSE_F = sum([(yi - yi_pred)**2 for yi, yi_pred in zip(actual, full)])
SSE_R = sum([(yi - yi_pred)**2 for yi, yi_pred in zip(actual, reduced)])
F_stat = ((SSE_R - SSE_F)/(dfr - dff))/(SSE_F/dff)
```

```

        return F_stat

def p_value()
    # F-test between model with all 5 features and model with 2 features
    dff, dfr = len(comp['actual'].values) - 5 - 1, len(comp['actual'].values) - 2 - 1
    F_stat = f_test(comp['predictions'].values, pred, comp['actual'].values, dff, dfr)
    F_stat

```

Out[83]: -1.7675408272236814

As far as we can see, we can't reject the null hypothesis. Which tell us that model with 2 components is better choice.

```

In [86]: # F-test between model with 2 features and null hypothesis
mean = sum(comp['actual'].values)/len(comp['actual'].values)
F_stat_null = f_test(pred, [mean]*40, comp['actual'].values, 37, 39)
F_stat_null

```

Out[86]: 4.918697454735734

But here, comparing with table value, we reject null hypothesis and confirm that model with two components performs better.

In conclusion, we choose linear model with **SIZELOG** and **CENTRAL** features and formula $y = \text{size}(-2.45) + \text{central}(-1.27) + 10.21$. R^2 is 0.21, which is not really good, but is the best choice among other linear models based on our research.

In []: